

A1: Data provenance and measurement audit

UCI Bank Marketing · Author: PLEASE FILL IN

This audit asks what the recorded fields can tell a dataset user before a marketing call. It examines the data and its documentation without fitting a prediction model.

操作：在 DS-A1 目录启动内核，按代码格 01 至 10 执行。七个报告部分覆盖全部必需审计。Data Card 和个人核验记录分别在 DATA_CARD.md 、 AI_USE_LOG.md 中维护，最后一格会显示它们。

1. Question and scope

Can dataset users treat all 20 recorded inputs as information available immediately before the recorded last marketing call? One documented post-call input is enough to reject this interpretation. Whether the remaining inputs are useful predictors is a separate question.

Concept	Definition for this audit
Stakeholders	Primarily researchers and data scientists reusing the table; also its providers and maintainers, bank decision makers, and affected customers.
Population	Records from the source bank's telephone marketing campaigns during the documented period. The public documentation does not fully establish coverage or campaign selection.
Sample	The 41,188 published records in the additional-full variant, documented as date ordered from May 2008 to November 2010.
Unit	A published client/campaign record summarizing the last contact and contact history. Without customer IDs, one row cannot be assumed to equal one unique customer or every individual call.
Target	y , recorded term-deposit subscription. The exact outcome timestamp and follow-up window are unspecified.

Concept	Definition for this audit
Estimand	Subscription probability conditional on information available before the last call, among the recorded campaigns. The audit checks whether the data support this quantity; it does not estimate that probability function.

The overall sample subscription fraction is descriptive. It does not estimate the conditional probability for a particular client. Also, the last call is identified retrospectively: a real application would need a decision time and outcome window specified in advance.

Code 01: setup

先确认内核工作目录为 `DS-A1` 。每次完整运行只增加一个时间戳结果目录，以保留旧结果和失败证据。

```
In [1]: from pathlib import Path
from datetime import datetime, timezone
import os, sys, json, hashlib, platform, random, subprocess

BASE = Path.cwd()
if not (BASE / "data/raw/manifest.json").is_file():
    raise FileNotFoundError("Start the notebook in the DS-A1 directory; see README.")
for name in ["tmp", "matplotlib"]:
    (BASE / ".runtime" / name).mkdir(parents=True, exist_ok=True)
os.environ["TMPDIR"] = str(BASE / ".runtime/tmp")
os.environ["MPLCONFIGDIR"] = str(BASE / ".runtime/matplotlib")
import tempfile
tempfile.tempdir = os.environ["TMPDIR"]
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from IPython.display import display, Markdown
from importlib.metadata import version

SEED = 20260928
random.seed(SEED)
np.random.seed(SEED)
RAW = BASE / "data/raw"
OUT = BASE / "results" /
datetime.now(timezone.utc).strftime("%Y%m%dT%H%M%S%fZ")
OUT.mkdir(parents=True)
versions = {p: version(p) for p in ["numpy", "pandas", "matplotlib",
"ipython", "ipykernel", "nbformat", "nbconvert"]}
environment = {"python": sys.version, "platform": platform.platform(), "seed": SEED, "packages": versions}
(OUT / "environment.json").write_text(json.dumps(environment, indent=2))
freeze = subprocess.run([sys.executable, "-m", "pip", "freeze"],
capture_output=True, text=True)
(OUT / "requirements-lock.txt").write_text(freeze.stdout)
if freeze.returncode:
    (OUT / "environment-error.txt").write_text(freeze.stderr)
    raise RuntimeError("Could not record dependencies; see environment-
```

```

error.txt.")
pd.set_option("display.max_rows", 25)
plt.rcParams.update({"figure.figsize": (8, 4), "font.size": 10})
print("Result directory:", OUT)
display(pd.Series(versions, name="version").to_frame())

```

Result directory: /data2/xuhongbo/DS_HW/DS-A1/results/20260929T014556_497511Z

	version
numpy	1.26.4
pandas	2.3.3
matplotlib	3.10.8
ipython	8.39.0
ipykernel	7.3.0
nbformat	5.11.1
nbconvert	7.17.1

2. Provenance and data-generating process

The data come from the [UCI repository](#). Attribution: Moro, S., Rita, P., & Cortez, P. (2014), *Bank Marketing*, DOI 10.24432/C5K306. UCI lists [CC BY 4.0](#) as the license. The local `bank-additional-names.txt` describes the selected variant. Legacy descriptions on the landing page use different encodings, including a different `pdays` sentinel.

Code 02: verify source bytes

The source URL, retrieval date, license, and SHA-256 values are recorded below. A hash mismatch stops the audit. Matching hashes establish unchanged bytes, not correct measurements.

```

In [2]: manifest = json.loads((RAW / "manifest.json").read_text())
hash_rows = []
for name, expected_hash in manifest["sha256"].items():
    actual = hashlib.sha256((RAW / name).read_bytes()).hexdigest()
    hash_rows.append({"file": name, "sha256": actual, "match": actual ==
expected_hash})
hashes = pd.DataFrame(hash_rows)
hashes.to_csv(OUT / "hash_verification.csv", index=False)
assert hashes["match"].all(), "Frozen data changed; investigate before
continuing."
(OUT / "source_manifest.json").write_text(json.dumps(manifest, indent=2))
print("Source:", manifest["source_url"])
print("Retrieved (UTC):", manifest["retrieved_utc"])
print("License:", manifest["license"])
for row in hash_rows:
    print(row["file"], "\n SHA-256:", row["sha256"], "\n match:",
row["match"])
docs = (RAW / "bank-additional-names.txt").read_text(encoding="utf-8-sig")
print("Local variant documentation loaded:", len(docs), "characters")

```


Source: <https://archive.ics.uci.edu/static/public/222/bank+marketing.zip>
Retrieved (UTC): 2026-09-28T08:32:02.794284+00:00
License: CC BY 4.0
bank-marketing.zip
SHA-256: e0bf5f5de5b846e2f18e9d90606637267d46dfa260e0f17bb12e605db5efbeb4
match: True
bank-additional-full.csv
SHA-256: 74adfc578bf77a7ff4bb1ba4a9f8709d9e3c6907342959c2c8416847e0afb4d8
match: True
bank-additional-names.txt
SHA-256: 10045462a4c026f858dfffa159370c4270510230c1fc63c6c91ca51bba4b3df26
match: True
Local variant documentation loaded: 5382 characters

Code 03: DGP and source lineage

The first row sketches how records arise; the second follows their path to this audit. The selection mechanism and the creators' processing are only partly documented. This is a process diagram, not an identified causal model.

```
In [3]: from matplotlib.patches import FancyBboxPatch
fig, axes = plt.subplots(2, 1, figsize=(9, 5))
flows = [
    ["Source bank clients\n(selection unknown)", "Selected
campaigns\nclient/history fields",
     "t0: pre-call\ndecision point", "Call occurs\nduration observed",
     "Recorded y\nhorizon unspecified"],
    ["Bank records +\nmacro indicators", "Creators' processing\nnot fully
observable",
     "UCI official archive\nadditional variant", "Frozen local bytes\nSHA-256
manifest",
     "Read-only audits\ntables / report"]]
for ax, labels, title in zip(axes, flows, ["Data-generating process
(schematic)", "Provenance lineage"]):
    ax.set(xlim=(0, 10), ylim=(0, 2))
    ax.axis("off")
    ax.set_title(title)
    for i, label in enumerate(labels):
        x = i * 2
        ax.add_patch(FancyBboxPatch((x + .05, .55), 1.75, .85,
                                     boxstyle="round,pad=0.03", facecolor="#e8f1fa",
                                     edgecolor="#496b88"))
        ax.text(x + .925, .975, label, ha="center", va="center", fontsize=8)
        if i < 4:
            ax.annotate("", xy=(x + 2.03, .975), xytext=(x + 1.82, .975),
                       arrowprops={"arrowstyle": "->"})
fig.tight_layout()
fig.savefig(OUT / "dgp_lineage.png", bbox_inches="tight")
plt.show()
```

Data-generating process (schematic)



Provenance lineage



3. Schema and data dictionary

A correct row count is not enough: the file must also contain the intended fields, types, and category encodings. `DATA_DICTIONARY.csv` records the 21 field definitions, units, domains, and timing judgments from the variant documentation. It is maintained separately from execution.

Code 04: read and check schema

The code preserves the raw values. Unexpected types or categories remain visible in the saved schema table; a wrong variant or column layout stops later analysis.

```

In [4]: dictionary = pd.read_csv(BASE / "DATA_DICTIONARY.csv", keep_default_na=False)
df = pd.read_csv(RAW / "bank-additional-full.csv", sep=";")
expected = dictionary.field.tolist()
if df.shape != (41188, 21) or df.columns.tolist() != expected:
    raise ValueError("Unexpected table shape or columns; check the selected variant.")
NUMERIC = dictionary.loc[dictionary.expected_type.ne("categorical"), "field"].tolist()
INTEGER = dictionary.loc[dictionary.expected_type.eq("integer"), "field"].tolist()
schema_rows = []
for row in dictionary.itertuples():
    s = df[row.field]
    numeric = pd.api.types.is_numeric_dtype(s)
    type_ok = numeric if row.expected_type != "categorical" else not numeric
    if row.expected_type == "integer":
        type_ok = type_ok and s.dropna().mod(1).eq(0).all()
    extra = sorted(set(s.dropna()) - set(row.documented_domain.split(" | ")))
    if row.expected_type == "categorical" else []
    schema_rows.append({"field": row.field, "dtype": str(s.dtype), "distinct": s.nunique(),
                        "type_ok": bool(type_ok), "unexpected_categories": ",
    ".join(extra)})
schema = pd.DataFrame(schema_rows)
  
```

```

schema.to_csv(OUT / "schema.csv", index=False)
display(schema.set_index("field"))
print("Sample subscription fraction:", round(df.y.eq("yes").mean(), 4))
print("Definitions and units:")
display(dictionary[["field", "meaning", "unit"]].set_index("field"))

```

	dtype	distinct	type_ok	unexpected_categories
field				
age	int64	78	True	
job	object	12	True	
marital	object	4	True	
education	object	8	True	
default	object	3	True	
housing	object	3	True	
loan	object	3	True	
contact	object	2	True	
month	object	10	True	
day_of_week	object	5	True	
duration	int64	1544	True	
campaign	int64	42	True	
pdays	int64	27	True	
previous	int64	8	True	
poutcome	object	3	True	
emp.var.rate	float64	10	True	
cons.price.idx	float64	26	True	
cons.conf.idx	float64	26	True	
euribor3m	float64	316	True	
nr.employed	float64	11	True	
y	object	2	True	

Sample subscription fraction: 0.1127

Definitions and units:

field	meaning	unit
age	Client age	years
job	Job category	category
marital	Marital status; divorced includes widowed	category
education	Education category	category
default	Credit in default	category
housing	Housing loan	category
loan	Personal loan	category
contact	Last contact channel	category
month	Month of last contact	month category
day_of_week	Weekday of last contact	weekday category
duration	Last call duration	seconds
campaign	Contacts in current campaign, including last	count
pdays	Days since prior-campaign contact; 999 denotes...	days or sentinel
previous	Contacts before current campaign	count
poutcome	Previous campaign outcome	category
emp.var.rate	Employment variation rate, quarterly	rate (scale not fully specified)
cons.price.idx	Consumer price index, monthly	index
cons.conf.idx	Consumer confidence index, monthly	index
euribor3m	Three-month Euribor, daily	rate
nr.employed	Number employed, quarterly	scale not specified in local readme
y	Term deposit subscribed	yes/no

4. Missingness and special codes

A parser can read `unknown` as ordinary text and report no nulls. That does not make the underlying measurement complete. The source documentation identifies this label as missing information, but does not explain its mechanism.

`pdays=999` has a different role: it denotes no previous contact. Treating it as a continuous value creates a false distance interpretation. Removing the sentinel


from a mean changes its scope to records with a recorded elapsed-day value; it does not recover a mean for all records.

Code 05: count and interpret encodings

```
In [5]: missing = pd.DataFrame({"parser_null_n": df.isna().sum(), "unknown_n":
df.eq("unknown").sum()})
missing["unknown_pct"] = missing.unknown_n / len(df) * 100
missing.to_csv(OUT / "missingness.csv", index_label="field")
display(missing)
unknown_any = df.eq("unknown").any(axis=1)
valid_pdays = df.loc[df.pdays.ne(999), "pdays"]
sentinel = pd.Series({"pdays_999_n": int(df.pdays.eq(999).sum()),
    "elapsed_day_records_n": len(valid_pdays), "mean_including_999":
df.pdays.mean(),
    "mean_excluding_999": valid_pdays.mean(), "rows_with_unknown":
int(unknown_any.sum())})
sentinel.to_csv(OUT / "sentinel_summary.csv", header=["value"])
display(sentinel.to_frame("value"))
fig, ax = plt.subplots()
missing.loc[missing.unknown_n.gt(0),
"unknown_pct"].sort_values().plot.barh(ax=ax)
ax.set(xlabel="Unknown-coded records (%)", ylabel="Field")
fig.tight_layout()
fig.savefig(OUT / "missingness.png", bbox_inches="tight")
plt.show()
print(f"The parser found {int(missing.parser_null_n.sum()):,} null cells; "
    f"the encoding check found {int(missing.unknown_n.sum()):,} unknown
entries.")
```


	parser_null_n	unknown_n	unknown_pct
age	0	0	0.000000
job	0	330	0.801204
marital	0	80	0.194231
education	0	1731	4.202680
default	0	8597	20.872584
housing	0	990	2.403613
loan	0	990	2.403613
contact	0	0	0.000000
month	0	0	0.000000
day_of_week	0	0	0.000000
duration	0	0	0.000000
campaign	0	0	0.000000
pdays	0	0	0.000000
previous	0	0	0.000000
poutcome	0	0	0.000000
emp.var.rate	0	0	0.000000
cons.price.idx	0	0	0.000000
cons.conf.idx	0	0	0.000000
euribor3m	0	0	0.000000
nr.employed	0	0	0.000000
y	0	0	0.000000

	value
pdays_999_n	39673.000000
elapsed_day_records_n	1515.000000
mean_including_999	962.475454
mean_excluding_999	6.014521
rows_with_unknown	10700.000000



The parser found 0 null cells; the encoding check found 12,718 unknown entries.

5. Range, duplicate, and anomaly checks

Range checks ask whether values fit their documented meaning. Negative employment variation or consumer confidence values can be meaningful. A blanket ban on negative numbers would therefore be inappropriate. Cross-field flags ask whether the contact-history encodings agree; any tension needs explanation before these fields are used together.

Code 06: ranges and contact-history consistency

```
In [6]: ranges = df[NUMERIC].describe(percentiles=[.01, .5, .99]).T
ranges.to_csv(OUT / "numeric_ranges.csv", index_label="field")
display(ranges[["min", "1%", "50%", "99%", "max"]].round(3))
lower_bounds = {"age": 1, "duration": 0, "campaign": 1, "pdays": 0,
"previous": 0}
range_flags = pd.DataFrame([{"field": col, "lower_bound":
lower_bounds.get(col, "not imposed"),
"below_bound_n": int(df[col].lt(lower_bounds[col]).sum()) if col in
lower_bounds else 0,
"nonfinite_n": int((~np.isfinite(df[col])).sum())} for col in NUMERIC])
range_flags.to_csv(OUT / "range_flags.csv", index=False)
display(range_flags)
cross_masks = {
"previous=0 but outcome exists": df.previous.eq(0) &
df.poutcome.ne("nonexistent"),
"previous>0 but outcome nonexistent": df.previous.gt(0) &
df.poutcome.eq("nonexistent"),
"previous=0 but elapsed days recorded": df.previous.eq(0) &
df.pdays.ne(999),
"previous>0 but pdays=999 (review)": df.previous.gt(0) & df.pdays.eq(999)}
cross = pd.Series({name: int(mask.sum()) for name, mask in
cross_masks.items()}, name="flagged_n")
cross.to_csv(OUT / "contact_history_flags.csv")
display(cross.to_frame())
print("These counts expose encoding tensions; they do not establish which
field is wrong.")
```

	min	1%	50%	99%	max
age	17.000	23.000	38.000	71.000	98.000
duration	0.000	11.000	180.000	1271.130	4918.000
campaign	1.000	1.000	2.000	14.000	56.000
pdays	0.000	3.000	999.000	999.000	999.000
previous	0.000	0.000	0.000	2.000	7.000
emp.var.rate	-3.400	-3.400	1.100	1.400	1.400
cons.price.idx	92.201	92.201	93.749	94.465	94.767
cons.conf.idx	-50.800	-49.500	-41.800	-26.900	-26.900
euribor3m	0.634	0.658	4.857	4.968	5.045
nr.employed	4963.600	4963.600	5191.000	5228.100	5228.100

	field	lower_bound	below_bound_n	nonfinite_n
0	age	1	0	0
1	duration	0	0	0
2	campaign	1	0	0
3	pdays	0	0	0
4	previous	0	0	0
5	emp.var.rate	not imposed	0	0
6	cons.price.idx	not imposed	0	0
7	cons.conf.idx	not imposed	0	0
8	euribor3m	not imposed	0	0
9	nr.employed	not imposed	0	0

	flagged_n
previous=0 but outcome exists	0
previous>0 but outcome nonexistent	0
previous=0 but elapsed days recorded	0
previous>0 but pdays=999 (review)	4110

These counts expose encoding tensions; they do not establish which field is wrong.

The frozen table contains 4,110 records with `previous > 0` and `pdays = 999`. Read literally, the documentation describes previous campaign contacts in one field and no previous contact in the other. This tension limits a simple joint

interpretation of the two fields. The public file does not resolve it, so the audit flags these records without deciding which field is wrong or deleting the rows.

Code 07: duplicates and a short anomaly review

Equal rows do not prove equal customer identities. The primary analysis retains them; a separate comparison reports what happens to the sample fraction if exact duplicates are removed.

The anomaly rules below are exploratory review thresholds, not official validity limits. Age above 100, a zero or hour-long call, and more than 20 campaign contacts deserve inspection. They can also be legitimate observations. The table is kept short because the important question is what produced a value, not how many outlier rules can flag it.

```
In [7]: XCOLS = [c for c in df if c != "y"]
exact_extra = int(df.duplicated().sum())
conflicts = int(df.groupby(XCOLS, dropna=False).y.nunique().gt(1).sum())
duplicates = pd.Series({"exact_extra_rows": exact_extra,
                        "rows_in_duplicate_groups": int(df.duplicated(keep=False).sum()),
                        "feature_only_extra_rows": int(df.duplicated(subset=XCOLS).sum()),
                        "feature_groups_with_conflicting_labels": conflicts}, name="count")
duplicates.to_csv(OUT / "duplicates.csv")
display(duplicates.to_frame())
df.loc[df.duplicated(keep=False)].to_csv(OUT / "duplicate_records.csv",
index_label="row_index")
dedup = df.drop_duplicates()
sensitivity = pd.DataFrame({"version": ["all rows", "deduplicated (sensitivity only)"],
                            "n": [len(df), len(dedup)], "subscription_fraction":
[df.y.eq("yes").mean(), dedup.y.eq("yes").mean()]})
sensitivity.to_csv(OUT / "dedup_sensitivity.csv", index=False)
display(sensitivity)
anomaly_masks = {"age > 100": df.age.gt(100), "duration = 0":
df.duration.eq(0),
                  "duration > 3600 seconds": df.duration.gt(3600), "campaign >
20": df.campaign.gt(20)}
anomalies = pd.DataFrame([{"review_rule": name, "n": int(mask.sum()),
                           "subscription_fraction": df.loc[mask, "y"].eq("yes").mean()} for name,
mask in anomaly_masks.items()])
anomalies.to_csv(OUT / "anomalies.csv", index=False)
display(anomalies)
print("All rows remain in the primary table. Empty subsets have undefined
rates, shown as NaN.")
```

	count
exact_extra_rows	12
rows_in_duplicate_groups	24
feature_only_extra_rows	12
feature_groups_with_conflicting_labels	0

	version	n	subscription_fraction
0	all rows	41188	0.112654
1	deduplicated (sensitivity only)	41176	0.112663

	review_rule	n	subscription_fraction
0	age > 100	0	NaN
1	duration = 0	4	0.000000
2	duration > 3600 seconds	5	0.600000
3	campaign > 20	157	0.006369

All rows remain in the primary table. Empty subsets have undefined rates, shown as NaN.

6. Leakage, counterexamples, and claim boundaries

The documentation states that `duration` is unavailable before the call. That timing is the reason it cannot be a predictor at the chosen decision point. The next table describes the association with `y`; it is not needed to prove the timing violation and is not a measure of predictive accuracy.

Code 08: availability and duration groups


Other predictors are marked conditional because this file does not establish their historical availability. Campaign counts need a decision-time definition, and macroeconomic indicators need publication dates and revision histories.

```
In [8]: availability = dictionary[["field", "availability", "timing_reason"]]
availability.to_csv(OUT / "feature_availability.csv", index=False)
display(availability.set_index("field"))
bins = [-1, 0, 60, 180, 300, 600, float("inf")]
labels = ["0", "1-60", "61-180", "181-300", "301-600", ">600"]
duration_bin = pd.cut(df.duration, bins=bins, labels=labels)
leakage = pd.DataFrame({"duration_seconds": duration_bin, "yes":
df.y.eq("yes").astype(int)})
leakage = leakage.groupby("duration_seconds", observed=False).yes.agg(["size",
"sum", "mean"])
leakage.columns = ["n", "yes_n", "subscription_fraction"]
leakage.to_csv(OUT / "duration_label_association.csv")
display(leakage)
fig, ax = plt.subplots()
leakage.subscription_fraction.plot.bar(ax=ax, rot=0, color="#437da8")
ax.set(xlabel="Recorded last-call duration (seconds)", ylabel="Sample
subscription fraction",
title="Post-call association: not usable as pre-call evidence")
fig.tight_layout(); fig.savefig(OUT / "duration_association.png",
```

```
bbox_inches="tight"); plt.show()
print("Availability verdict: reject treating all 20 inputs as pre-call
information.")
print("Dropping duration is necessary, but does not verify remaining field
timestamps.")
```

field	availability	timing_reason
age	conditional	Client attribute; historical snapshot unverified
job	conditional	Client attribute; historical snapshot unverified
marital	conditional	Client attribute; historical snapshot unverified
education	conditional	Client attribute; historical snapshot unverified
default	conditional	Client attribute; historical snapshot unverified
housing	conditional	Client attribute; historical snapshot unverified
loan	conditional	Client attribute; historical snapshot unverified
contact	conditional	Must be planned and available before this call
month	conditional	Calendar is known; last-call selection is retr...
day_of_week	conditional	Calendar is known; last-call selection is retr...
duration	exclude	Post-call: unavailable at t0
campaign	conditional	Requires decision-time count; final total may ...
pdays	conditional	History candidate; validate snapshot and senti...
previous	conditional	History candidate; validate snapshot
poutcome	conditional	History candidate; validate snapshot
emp.var.rate	conditional	Publication lag and revisions require as-of re...
cons.price.idx	conditional	Publication lag and revisions require as-of re...
cons.conf.idx	conditional	Publication lag and revisions require as-of re...
euribor3m	conditional	Daily value must be published before t0
nr.employed	conditional	Publication lag and revisions require as-of re...
y	target	Outcome only; never a predictor

	n	yes_n	subscription_fraction
duration_seconds			
0	4	0	0.000000
1-60	4282	1	0.000234
61-180	16419	563	0.034290
181-300	9279	954	0.102813
301-600	7740	1438	0.185788
>600	3464	1684	0.486143



Availability verdict: reject treating all 20 inputs as pre-call information. Dropping duration is necessary, but does not verify remaining field timestamps.

Interpretation

`duration` is a counterexample to the claim that all provided inputs can be used before a call. Deleting it is necessary for that task, but other fields still need availability checks. The record of the last call also comes from a retrospective view of a campaign.

The missingness audit provides a separate measurement counterexample: zero parser nulls can coexist with unknown-coded observations. Similarly, averaging `pdays` with 999 included mixes an elapsed-time measurement with a contact-status code.

A selection counterexample is also possible. If the bank preferentially contacted customers already likely to subscribe, the sample fraction could differ from that among uncontacted customers. The table does not identify whether or how strongly this selection occurred.

The following claims are not supported:

- Using all provided features demonstrates prediction before calling.

- No NaN values means that all measurements are complete.
- The recorded subscription fraction applies to all current bank customers.
- Longer calls cause more subscriptions.

The audit establishes a documented timing violation and describes measurement problems in this frozen table. It does not establish model accuracy or causal effects. The lack of customer IDs prevents reliable customer-level split checks. Documented date ordering, without complete timestamps, is insufficient to validate a deployment-style temporal evaluation.

7. Findings, Data Card, and reflection

Code 09: save the numerical findings

The figures below come from this execution. Earlier output directories remain available, including incomplete runs. The source hashes and dictionary are retained with this run's record so its provenance can be checked later.

```
In [9]: findings = {"rows": len(df), "columns": len(df.columns),
    "subscription_fraction": float(df.y.eq("yes").mean()),
    "parser_null_cells": int(missing.parser_null_n.sum()),
    "unknown_cells": int(missing.unknown_n.sum()), "rows_with_unknown":
int(unknown_any.sum()),
    "pdays_999_rows": int(df.pdays.eq(999).sum()), "exact_extra_rows":
exact_extra,
    "conflicting_feature_groups": conflicts,
    "schema_type_flags": int((~schema.type_ok).sum()),
    "schema_category_flags": int(schema.unexpected_categories.ne("").sum()),
    "range_flags": int(range_flags.below_bound_n.sum() +
range_flags.nonfinite_n.sum()),
    "availability_verdict": "Reject all-input pre-call interpretation:
duration is observed after the call"}
(OUT / "findings.json").write_text(json.dumps(findings, indent=2))
display(pd.Series(findings, name="result").to_frame())
for name in ["DATA_DICTIONARY.csv", "requirements.txt"]:
    (OUT / name).write_bytes((BASE / name).read_bytes())
print("Review contact-history and anomaly flags even when schema/range checks
pass.")
```

	result
rows	41188
columns	21
subscription_fraction	0.112654
parser_null_cells	0
unknown_cells	12718
rows_with_unknown	10700
pdays_999_rows	39673
exact_extra_rows	12
conflicting_feature_groups	0
schema_type_flags	0
schema_category_flags	0
range_flags	0
availability_verdict	Reject all-input pre-call interpretation: dura...

Review contact-history and anomaly flags even when schema/range checks pass.

Code 10: display the maintained report files

Edit `DATA_CARD.md` and complete the student section of `AI_USE_LOG.md`, then rerun this cell and save the notebook. The assistant's source review is separate from your own verification. `SUBMISSION_SUMMARY.md` is a result-summary draft for review before platform submission.

```
In [10]: for name in ["DATA_CARD.md", "AI_USE_LOG.md", "SUBMISSION_SUMMARY.md"]:
        report = (BASE / name).read_text()
        display(Markdown(report))
        (OUT / name).write_text(report)
        print("Saved run:", OUT)
        print("Complete personal verification, save A1.ipynb, then run
        export_report.py.")
```

Data card: UCI Bank Marketing

This card describes the `bank-additional-full.csv` variant. It is maintained as report text; executing the notebook does not rewrite it. Numerical audit findings are in the saved notebook and its `results` directory. Student review is pending.

Source and permitted use

Moro, S., Rita, P., & Cortez, P. (2014). *Bank Marketing*. UCI Machine Learning Repository. [Dataset DOI](#). The [UCI page](#) lists [CC BY 4.0](#), which requires attribution. Audit tables and plots are derived artifacts; the raw CSV has not been modified.

- Download: <https://archive.ics.uci.edu/static/public/222/bank+marketing.zip>
- Retrieved UTC: 2026-09-28T08:32:02.794284+00:00
- CSV SHA-256:
`74adfc578bf77a7ff4bb1ba4a9f8709d9e3c6907342959c2c8416847e0afb4d8`
- Full archive and documentation hashes: `data/raw/manifest.json`.
- Variant definitions: `data/raw/bank-additional-names.txt`.
- Related paper: Moro, Cortez & Rita (2014), DOI 10.1016/j.dss.2014.03.001. This audit does not replicate the paper.

Records and question

The documentation describes 41,188 records with 20 inputs and the target `y`, ordered by date from May 2008 to November 2010. A published row summarizes a client/campaign record, including the last contact and earlier contacts. The file has no customer identifier, so a row is not assumed to be a unique person or a record of every call.

The audit asks whether every input was available immediately before the recorded last call. Its main stakeholders are researchers reusing the table. Providers, maintainers, bank decision makers, and affected customers also have an interest in valid interpretation. The intended population is records from the source bank's telephone marketing campaigns in this period; its coverage and selection mechanism are incompletely documented. The quantity of interest is subscription probability conditional on information available before the call, among recorded campaigns. This audit checks support for that quantity; it does not estimate the conditional function or train a model.

Measurement and handling

`y` records term-deposit subscription; the exact follow-up window is unspecified. `unknown` marks unrecorded/unknown categories. `pdays=999` is a status code for no previous contact, not 999 elapsed days. The main analysis preserves all rows and

original values. Duplicate and extreme-value flags prompt review, not automatic deletion. A separate exact-deduplication comparison shows sensitivity without changing the main table. Field definitions, units, domains, and timing judgments are in `DATA_DICTIONARY.csv` .

Limits

`duration` is observed after the call and cannot support prediction before it. The remaining inputs require historical snapshot or publication-time evidence, especially campaign totals and macroeconomic indicators. Removing `duration` alone does not establish that the predictors are valid at the decision time. Missing customer IDs prevent reliable customer-level deduplication and split checks. The public table does not establish representativeness beyond the recorded campaigns. It supports descriptive auditing, not claims of deployment accuracy or causal effects of calling. No identifiers are added and no person-level linkage is performed.

Reproduction and review

Run `A1.ipynb` from the repository directory using the dependencies in `requirements.txt` . Each run records its environment and data hashes. The notebook displays this card in its PDF output; update this file and rerun that display cell if its text changes. AI assistance and the student's separate verification record are in `AI_USE_LOG.md` .

AI use and independent verification

This draft was prepared with ChatGPT and Codex. Codex read the assignment PDF and the variant-specific source documentation, wrote and tested the audit code, and revised the structure after the supplied ChatGPT review. These are actions by the assistant. They are not evidence that the student has independently checked the work.

Suggestion	Decision in this draft	Reason or evidence
Use Bank Marketing and ask about information available before a call	Accepted	The additional-variant documentation identifies a post-call variable.
Separate <code>unknown</code> from parser nulls and <code>pdays=999</code> from elapsed days	Accepted	Attribute 13 and section 8 of the local names file document these encodings.
Shorten the notebook and maintain the Data Card separately	Accepted	The same required audits fit into seven report sections.
Replace the estimand with a proportion/probability without conditioning	Modified	Plain language is clearer, but the sample fraction and conditional probability are different quantities.
Define population as whatever the dataset claims to represent	Modified	The source setting is stated explicitly; actual coverage and selection remain uncertain.
Use <code>Path.cwd()</code> as the project root	Accepted with a directory check	Execution now requires starting in DS-A1, as documented in README.
Treat extreme values as errors	Not adopted	The review thresholds are exploratory; they do not establish recording errors.

Student verification and reflection (complete personally)

- Name and date: TODO
- Source definition I checked, including file/section and finding: TODO
- Numerical finding I checked independently, including method and result: TODO
- AI advice I accepted or rejected and my reason: TODO
- Unexpected or negative evidence and its effect on my conclusion: TODO
- My answer to the audit question and its limits: TODO

Read the original documentation and use a separate calculation, spreadsheet, or manual raw-row check for at least one numerical result. Record what you actually

did. Then rerun the notebook's last cell so the saved output and PDF include the revised log.

Submission summary draft

This audit asks whether users of the UCI Bank Marketing additional-full dataset can treat all 20 recorded inputs as information available before the last marketing call. It examines 41,188 records and preserves the source attribution, license, retrieval date, and SHA-256 hashes. The intended quantity is subscription probability conditional on information available before a call, within the recorded campaigns. No model is trained.

The checks cover schema, ranges, missingness, duplicates, anomalies, and information availability. Although the table has no parser-null cells, it contains 12,718 unknown-coded entries across 10,700 records. The pdays sentinel occurs in 39,673 rows and cannot be interpreted as elapsed days. There are 12 excess identical rows. They remain in the primary analysis because identical values do not establish customer identity. The observed subscription fraction is about 11.27%, which describes this sample.

The main counterexample is duration, which the documentation places after the call. Its association with subscription cannot justify prediction before calling. Removing duration is necessary but does not verify the availability of campaign totals, historical customer attributes, or macroeconomic releases. Campaign selection, missing customer identifiers, and unspecified outcome timing also limit interpretation. The audit therefore rejects an all-input pre-call interpretation while making no claim about predictive accuracy, causality, or deployment readiness.

Repository URL and commit/tag: TODO

Student review and independent verification: complete AI_USE_LOG.md before submission.

Saved run: /data2/xuhongbo/DS_HW/DS-A1/results/20260929T014556_497511Z
Complete personal verification, save A1.ipynb, then run export_report.py.